



# ARM CORTEX ASSEMBLER PRIMER



**KEVIN THOMAS**  
**DIRECTOR OF TEST AUTOMATION**

[mytechnotalent.com](http://mytechnotalent.com)



# AGENDA

Linker Script	3	MOVW & MOVT	10
Vector Table	4	Flags	11
Registers	5	Bit Set & Clear	12
MOV Immediate	6	Conditional Branching	13
MOVS Register	7	Loops	14
LDR Immediate	8	Interrupts	15
LDR Offset	9	STUXNET Demo	16

# LINKER SCRIPT

On p50 of the STM32F401CCUx datasheet, we see the memory map. We notice that starting at 0x00000000 we have flash aliased which will be mapped to 0x08000000, flash memory during the linking process. We also see SRAM starting at 0x20000000 and we see our peripherals starting at 0x40000000. On p43 of the reference manual we see another table indicating similar info. On page 45 of the reference manual we see how the flash sectors are organized. We begin by creating our linker script, st\_nucleo\_f4.ld and define the MEMORY and SECTIONS based on p50 of the STM32F401CCUx datasheet. When we write Assembler instructions follow three steps.

1. assemble (creates a relocatable object file)  
arm-none-eabi-as -g main.s -o main.o
2. link (map all the 0x0000000 addresses to physical addresses)  
arm-none-eabi-ld main.o -o main.elf -T st\_nucleo\_f4.ld
3. flash  
openocd -f interface/stlink-v2.cfg -f target/stm32f4x.cfg -c "program main.elf verify reset exit"

We can dive into the object file with the below commands.

- \* display the contents of the section headers  
arm-none-eabi-objdump -h main.o
- \* display assembler contents of all sections  
arm-none-eabi-objdump -D main.o | less
- \* display assembler contents of executable sections  
arm-none-eabi-objdump -d main.o
- \* display the full contents of all sections requested  
arm-none-eabi-objdump -s main.o | less
- \* display information about the contents of ELF format files  
arm-none-eabi-readelf -a main.elf | less
- \* list symbols in file  
arm-none-eabi-nm main.elf | less

Keep in mind when you are using objdump there may be data in the binary which objdump will try to interpret as instructions so you can ignore those parts.

```
/*
 * FILE: st_nucleo_f4.ld
 *
 * DESCRIPTION:
 * This file contains the linker script
 * utilizing the STM32F401 Nucleo-64 microcontroller.
 *
 * AUTHOR: Kevin Thomas
 * CREATION DATE: September 1, 2023
 * UPDATE Date: September 1, 2023
 */

MEMORY
{
    FLASH : ORIGIN = 0x08000000, LENGTH = 512K
    SRAM : ORIGIN = 0x20000000, LENGTH = 96K
}

SECTIONS
{
    .isr_vector :
    {
        *(.isr_vector)
    } >FLASH
    .text :
    {
        *(.text)
    } >FLASH
    .data (NOLOAD) :
    {
        . = . + 0x400;
        _estack = .;
        *(.data)
    } >SRAM
}
```

# VECTOR TABLE

The ARM Cortex M4 has a vector table which is the first thing loaded into memory on boot. On page 202 of the reference manual, we see the vector table. These will also get mapped during the linking process. The very base of memory is the MSP or master stack pointer followed by the Reset\_Handler function. The rest of the handlers follow the Reset\_Handler followed by the IRQ or interrupt requests.

- \* 85 IRQ's
- \* 15 system exceptions
- \* master stack pointer
- \*  $85+15+1 = 101 * 4 = 404$  bytes or 0x194 bytes

```
/*
 * The STM32F401CCUx vector table. Note that the proper constructs
 * must be placed on this to ensure that it ends up at physical address
 * 0x00000000.
 */
.global isr_vector
.section .isr_vector, "a"
.type isr_vector, %object
isr_vector:
.word _estack
.word Reset_Handler
weak NMI_Handler
weak HardFault_Handler
weak MemManage_Handler
weak BusFault_Handler
weak UsageFault_Handler
.word 0
.word 0
.word 0
.word 0
weak SVC_Handler
weak DebugMon_Handler
.word 0
weak PendSV_Handler
weak SysTick_Handler
.word 0
weak EXTI16_PVD_IRQHandler // EXTI Line 16 interrupt /PVD through EXTI line detection
interrupt
weak TAMP_STAMP_IRQHandler // Tamper andTimeStamp interrupts through the EXTI line
weak EXTI22_RTC_WKUP_IRQHandler // EXTI Line 22 interrupt /RTC Wakeup interrupt through the
EXTI line
weak FLASH_IRQHandler // FLASH global interrupt
weak RCC_IRQHandler // RCC global interrupt
...
```

# REGISTERS

The ARM Cortex M4 has 13 general-purpose registers (R0-R12), 1 stack pointer register (R13 or SP), 1 link register (R14 or LR), 1 program counter (R15 or PC) and a status register (xPSR) and 3 interrupt mask registers. There are peripheral registers as well located on page 38 of the reference manual. All registers are 32-bits wide or 4 bytes wide.

## \* ASSEMBLE AND LINK w/ SYMBOLS:

- \* 1. arm-none-eabi-as -g main.s -o main.o
- \* 2. arm-none-eabi-ld main.o -o main.elf -T st\_nucleo\_f4.ld
- \* 3. openocd -f interface/stlink-v2.cfg -f target/stm32f4x.cfg -c "program main.elf verify reset exit"

## \* DEBUG w/ SYMBOLS:

- \* 1. openocd -f board/st\_nucleo\_f4.cfg
- \* 2. arm-none-eabi-gdb main.elf
- \* 3. target remote :3333
- \* 4. monitor reset halt
- \* 5. b \_\_start
- \* 6. c

r0	0x40020400	1073873920
r1	0x282	642
r2	0x35010041	889258049
r3	0x0	0
r4	0x0	0
r5	0x0	0
r6	0x0	0
r7	0x42	66
r8	0x0	0
r9	0x0	0
r10	0x0	0
r11	0x0	0
r12	0x0	0
sp	0x20000400	0x20000400
lr	0x800022d	134218285
pc	0x800043c	0x800043c <Conditional_Branching>
xPSR	0x61000000	1627389952
fpscr	0x0	0
msp	0x20000400	0x20000400
psp	0x0	0x0
primask	0x0	0

# MOV IMMEDIATE

```
// Move an immediate value into R0.  
MOV_Immediate:  
    MOV R0, #0x42          // move 0x42 into R0  
    BX LR                  // return to caller
```

```
(gdb) disas  
Dump of assembler code for function MOV_Immediate:  
=> 0x080003ca <+0>:  mov.w  r0, #66 ; 0x42  
  0x080003ce <+4>:  bx    lr
```

```
(gdb) s  
halted: PC: 0x080003ce  
(gdb) x/x $r0  
0x42: 0x019d0000
```

# MOVS REGISTER

```
// Move an immediate value into R0 and R1 and then move R1 into R0.
```

```
MOVS_Register:  
MOV R0, #0x42      // move 0x42 into R0  
MOV R1, #0x43      // move 0x43 into R1  
MOVS R0, R1        // move R1 into R0  
BX LR              // return to caller
```

```
(gdb) disas  
Dump of assembler code for function MOVS_Register:  
=> 0x080003d0 <+0>:  mov.w  r0, #66 ; 0x42  
0x080003d4 <+4>:  mov.w  r1, #67 ; 0x43  
0x080003d8 <+8>:  movs   r0, r1  
0x080003da <+10>: bx    lr
```

```
(gdb) s  
(gdb) x/x $r0  
0x42: 0x019d0000
```

```
(gdb) s  
(gdb) x/x $r1  
0x43: 0x00019d00
```

```
(gdb) s  
(gdb) x/x $r0  
0x43: 0x00019d00
```

# LDR IMMEDIATE

/// Load an address into R0, load value inside R0 into R1, set a bit and then store back the value into the address.

```
LDR_Immediate:  
LDR R0, =0x40023830      // load address of RCC_AHB1ENR register (p118 RM)  
LDR R1, [R0]                // load value inside RCC_AHB1ENR register  
ORR R1, #(1<<1)          // set the GPIOBEN bit  
STR R1, [R0]                // store value into RCC_AHB1ENR register  
BX LR                      // return to caller
```

(gdb) disas

Dump of assembler code for function LDR\_Immediate:

```
=> 0x080003dc <+0>: ldr r0, [pc, #136] ; (0x8000468 <Interrupts+6>)  
0x080003de <+2>: ldr r1, [r0, #0]  
0x080003e0 <+4>: orr.w r1, r1, #2  
0x080003e4 <+8>: str r1, [r0, #0]  
0x080003e6 <+10>: bx lr
```

(gdb) s

(gdb) x/x \$r0

```
0x40023830: 0x00000007
```

(gdb) s

(gdb) s

(gdb) x/x \$r1

```
0x7: 0x00019d08
```

(gdb) s

(gdb) x/x \$r0

```
0x40023830: 0x00000007
```

# LDR OFFSET

```
// Load the SRAM base address into R0, load an offset address into R1 and left shift the offset in R1 by 1  
which multiplies the value  
// by 2 and then add it to the base value producing 0x20000020 to be stored back into R2. Set a 1 in the  
LSB and store back the value  
// into R0.  
LDR_Offset:  
LDR R0, =0x20000000          // SRAM base (p42 RM)  
LDR R1, =0x10                // offset address  
LDR R2, [R0, R1, LSL #1]      // 0x20000000 + (0x10 * 2) = 0x20000020, then store what is inside  
0x20000020 into R2  
ORR R2, #(1<<0)             // put a 1 in LSB with the contents inside R2  
STR R2, [R0]                  // store the contents inside R2 into 0x20000000  
BX LR                         // return to caller
```

```
(gdb) disas  
Dump of assembler code for function LDR_Offset:  
=> 0x080003e8 <+0>:  mov.w  r0, #536870912 ; 0x20000000  
0x080003ec <+4>:  mov.w  r1, #16  
0x080003f0 <+8>:  ldr.w  r2, [r0, r1, lsl #1]  
0x080003f4 <+12>: orr.w  r2, r2, #1  
0x080003f8 <+16>: str   r2, [r0, #0]  
0x080003fa <+18>: bx    lr
```

```
(gdb) s  
(gdb) x/x $r0  
0x20000000: 0x35010041  
  
(gdb) s  
(gdb) s  
(gdb) x/x $r2  
0x35010040: 0x00000000  
(gdb) x/x 0x20000020  
0x20000020: 0x35010040
```

```
(gdb) s  
(gdb) s  
(gdb) x/x $r2  
0x35010041: 0x00000000  
(gdb) x/x $r0  
0x20000000: 0x35010041
```

# MOVW MOVT

/// Move half word into the MSB of R0 and then another half word into the LSB of R0 and load an address in RAM and store that value into it.

```
MOVW_MOVT:  
    MOVW R0, #0x5678      // move half word into MSB of R0  
    MOVT R0, #0x1234      // move half word into LSB of R0  
    LDR  R1, =0x20002000  // load address in RAM  
    STR  R0, [R1]         // store word into the value stored inside R1  
    BX   LR               // return to caller
```

```
(gdb) disas  
Dump of assembler code for function MOVW_MOVT:  
=> 0x080003fc <+0>:  movw  r0, #22136   ; 0x5678  
     0x08000400 <+4>:  movt  r0, #4660    ; 0x1234  
     0x08000404 <+8>:  mov.w r1, #536879104 ; 0x20002000  
     0x08000408 <+12>: str   r0, [r1, #0]  
     0x0800040a <+14>: bx    lr
```

```
(gdb) s  
(gdb) x/x $r0  
0x5678: 0x00000000
```

```
(gdb) s  
(gdb) x/x $r0  
0x12345678: 0x00000000
```

```
(gdb) s  
(gdb) s  
(gdb) x/x $r1  
0x20002000: 0x12345678
```

# FLAGS

```
/// Move a negative value into R0 and R1 and add them to set the N, C and T bits in the xPSR  
(0b1010000100000000000000000000000000000000).  
// Load the max 32-bit signed value into R0 and load a value into R1 and add them to set the N, V and T  
bits inside the xPSR  
// (0b10010001000000000000000000000000). Load the max 32-bit unsigned value into R0 and load a value  
into R1 and add them to  
// set the Z, C and T bits inside the xPSR (0b01100001000000000000000000000000).  
Flags:  
MOV R0, #-1          // move a negative value into R0  
MOV R1, #-2          // move a negative value into R1  
ADDS R0, R1          // N, C, T bits set in xPSR (negative, carry and thumb) (p573 ARM)  
LDR R0, =0xFFFFFFFF  // load the max 32-bit signed value into R0  
LDR R1, =0x01         // load a value into R1 that when added to R0 will overflow  
ADDS R0, R1          // N, V, T bits set in xPSR (negative, overflow and thumb)  
LDR R0, =0xFFFFFFFF  // load the max 32-bit unsigned value into R0  
LDR R1, =0x01         // load a value into R1 that when added to R0 will overflow  
ADDS R0, R1          // Z, C, T bits set in xPSR (zero, carry, thumb)  
BX LR                // return to caller
```

```
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) x/x $xPSR  
0x91000000 // 0b100100010000000000000000000000000000000000000000
```

```
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) x/x $xPSR  
0x61000000 // 0b011000010000000000000000000000000000000000000000
```

# BIT SET CLEAR

```
// Load the GPIOB_MODER register into R0 and move the value inside the register into R1, perform at bit  
clear and then a bit set and store the value back into the register.
```

```
Bit_Set_Clear:  
LDR R0, =0x40020400      // load address of GPIOB_MODER register (p158 RM)  
LDR R1, [R0]                // load value inside GPIOB_MODER register  
AND R1, #~(1<<1)          // clear the MODER0 bit  
ORR R1, #(1<<1)           // set the MODER0 bit  
STR R1, [R0]                // store value into GPIOB_MODER register  
BX LR                      // return to caller
```

```
(gdb) disas  
Dump of assembler code for function Bit_Set_Clear:  
=> 0x0800042c <+0>: ldr r0, [pc, #112] ; (0x80004a0 <Interrupts+60>) 0x0800042e <+2>: ldr r1, [r0, #0]  
0x08000430 <+4>: bic.w r1, r1, #2  
0x08000434 <+8>: orr.w r1, r1, #2  
0x08000438 <+12>: str r1, [r0, #0]  
0x0800043a <+14>: bx lr
```

```
(gdb) s  
(gdb) s  
(gdb) x/x $r1  
0x282: 0x6801487c
```

```
(gdb) s  
(gdb) x/x $r1  
0x280: 0x487c6001
```

```
(gdb) s  
(gdb) x/x $r1  
0x282: 0x6801487c
```

```
(gdb) s  
(gdb) x/x $r0  
0x40020400: 0x00000282
```

# CONDITIONAL BRANCHING

```
// Load the GPIOB_MODER register into R0 and move the value inside the register into R1, perform at bit  
// clear and then a bit set and store the value back into the register.  
  
Bit_Set_Clear:  
    LDR R0, =0x40020400          // load address of GPIOB_MODER register (p158 RM)  
    LDR R1, [R0]                 // load value inside GPIOB_MODER register  
    AND R1, #~(1<<1)           // clear the MODER0 bit  
    ORR R1, #(1<<1)            // set the MODER0 bit// Load the EXTI_PR register into R0 and the  
contents into R1 and test if bit 13 is a 0 (0b00100000000000001000000000)  
  
// then move a random mock value into R0 and compare against another random mock value and branch  
appropriately.  
  
Conditional_Branching:  
    LDR R0, =0x40013C14 // load the address of EXTI_PR register (p211 RM)  
    LDR R1, [R0] // load value inside EXTI_PR register  
    TST R1, #(1<<13) // read the PR13 bit, if 0, then BEQ  
    BEQ .Conditional_TST_Finish // branch if equal  
    NOP // dummy logic  
.Conditional_TST_Finish:  
    MOV R0, #0x40 // random mock value  
    CMP R0, #0x40 // compare random mock value to another value  
    BNE .Conditional_TST_Finish // branch if not equal  
    BLE .Conditional_CMP_Finish // branch if less than or equal  
    NOP // dummy logic  
.Conditional_CMP_Finish:  
    BX LR // return to caller
```

```
(gdb) disas  
Dump of assembler code for function Conditional_Branching:  
=> 0x0800043c <+0>: ldr r0, [pc, #100] ; (0x80004a4 <Interrupts+64>  
0x0800043e <+2>: ldr r1, [r0, #0]  
0x08000440 <+4>: tst.w r1, #8192 ; 0x2000  
0x08000444 <+8>: beq.n 0x8000448 <.Conditional_TST_Finish>  
0x08000446 <+10>: nop  
End of assembler dump.
```

```
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) disas  
Dump of assembler code for function .Conditional_TST_Finish:  
=> 0x08000448 <+0>: mov.w r0, #64 ; 0x40  
0x0800044c <+4>: cmp r0, #64 ; 0x40  
0x0800044e <+6>: bne.n 0x8000448 <.Conditional_TST_Finish>  
0x08000450 <+8>: ble.n 0x8000454 <.Conditional_CMP_Finish>  
0x08000452 <+10>: nop
```

```
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) disas  
Dump of assembler code for function .Conditional_TST_Finish:  
0x08000448 <+0>: mov.w r0, #64 ; 0x40  
0x0800044c <+4>: cmp r0, #64 ; 0x40  
0x0800044e <+6>: bne.n 0x8000448 <.Conditional_TST_Finish>  
=> 0x08000450 <+8>: ble.n 0x8000454 <.Conditional_CMP_Finish>  
0x08000452 <+10>: nop
```

# LOOPS

```
// Move a value into R0 and we have a NOP representing any logic you would want run then subtract 1  
from the counter  
// and compare against 0 otherwise loop.  
Loops:  
MOV R0, #0x10          // move a value into R0  
.For_Loop:  
NOP                  // dummy logic  
SUBS R0, #0x01         // subtract a value from R0  
CMP R0, #0x00          // subtract R0 from a value without updating R0  
BNE .For_Loop         // branch if not equal  
BX LR                // return to caller
```

```
(gdb) disas  
Dump of assembler code for function Loops:  
=> 0x08000456 <+0>:  mov.w  r0, #16  
  
(gdb) s  
(gdb) disas  
Dump of assembler code for function .For_Loop:  
=> 0x0800045a <+0>:  nop  
 0x0800045c <+2>:  subs   r0, #1  
 0x0800045e <+4>:  cmp    r0, #0  
 0x08000460 <+6>:  bne.n  0x800045a <.For_Loop>  
 0x08000462 <+8>:  bx     lr  
  
(gdb) x/x $r0  
0x10: 0x0800019d
```

```
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) x/x $r0  
0x10: 0x0800019d  
(gdb) x/x $r0  
0xf: 0x00019d0
```

```
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) s  
(gdb) x/x $r0  
0x10: 0x0800019d  
(gdb) x/x $r0  
0xe: 0x019d0800  
...
```

# INTERRUPTS

```
// Review EXTI15_10_IRQHandler and USART1_IRQHandler. The vector table's weak reference must be  
// changed to a .word as we define the handler. Review blue button press and observe 'A' displayed  
// on UART and typing a char in the UART displays the char. These actions are independent of the  
// main loop and will interrupt the main loop.
```

Interrupts:

```
BX LR
```

```
/**  
 * @brief USART1 interrupt handler.  
 *  
 * @details This is the interrupt handler function for USART1 interrupts. It is called when  
 * USART1 generates an interrupt request. Inside this handler, it calls the  
 * `USART1_Callback` function to handle the USART1 interrupt, and then returns to  
 * the main program execution.  
 *  
 * @param None  
 * @retval None  
 */  
.section .text.USART1_IRQHandler  
.weak USART1_IRQHandler  
.type USART1_IRQHandler, %function  
USART1_IRQHandler:  
    PUSH {LR}          // push return address onto stack  
    LDR R0, =0x40011000 // load address of USART1_SR register  
    LDR R1, [R0]        // load value inside USART1_SR register  
    TST R1, #(1<<5)   // read the RXNE register  
    BL USART1_Callback // call function  
    POP {LR}          // pop return address from stack  
    BX LR             // return from the function
```

# STUXNET DEMO

In this scenario, we have been contracted by a classified organization to infiltrate the Natanz Nuclear Facility. The intel we have been given is that a normal rate of delay is 64 ms for the centrifuge to spin is based in information gathered from a classified source and that if it ever reached 8 it would spin the centrifuge out of control and destroy the facility which is our goal ;) The intel provided also explains that by entering in a 1 into the terminal will act as a kill switch in the event that an Operator observes something wrong.

The other piece of intel is that the motor is controlled by a UNL2003 driver with a stepper motor and it is designed in a full drive sequence mode which two coils are energized at a time that means two windings of stepper motor energized together. Therefore, motor runs at full torque.

Let's assemble...

```
arm-none-eabi-as -g main.s -o main.o  
arm-none-eabi-ld main.o -o main.elf -T st_nucleo_f4.ld  
arm-none-eabi-objcopy -O binary --strip-all main.elf main.bin  
openocd -f interface/stlink-v2.cfg -f target/stm32f4x.cfg -c "program main.bin 0x08000000 verify reset  
exit"
```

Terminal 1:

```
openocd -f board/st_nucleo_f4.cfg
```

Terminal 2:

```
arm-none-eabi-gdb main.bin  
target remote :3333  
monitor reset halt
```

# STUXNET DEMO

In this scenario, we have been contracted by a classified organization to infiltrate the Natanz Nuclear Facility. The intel we have been given is that a normal rate of delay is 64 ms for the centrifuge to spin is based in information gathered from a classified source and that if it ever reached 8 it would spin the centrifuge out of control and destroy the facility which is our goal ;) The intel provided also explains that by entering in a 1 into the terminal will act as a kill switch in the event that an Operator observes something wrong.

The other piece of intel is that the motor is controlled by a UNL2003 driver with a stepper motor and it is designed in a full drive sequence mode which two coils are energized at a time that means two windings of stepper motor energized together. Therefore, motor runs at full torque.

Let's assemble...

```
arm-none-eabi-as -g main.s -o main.o  
arm-none-eabi-ld main.o -o main.elf -T st_nucleo_f4.ld  
arm-none-eabi-objcopy -O binary --strip-all main.elf main.bin  
openocd -f interface/stlink-v2.cfg -f target/stm32f4x.cfg -c "program main.bin 0x08000000 verify reset  
exit"
```

Terminal 1:

```
openocd -f board/st_nucleo_f4.cfg
```

Terminal 2:

```
arm-none-eabi-gdb main.bin  
target remote :3333  
monitor reset halt
```

Let's look at first 1000 instructions.

```
(gdb) x/1000i 0x08000000
```

Arch Ref Manual page 682, we see NVIC\_ISER0-NVIC\_ISER15 so we see the address starting at 0xE000E100 to 0xE000E13C. So lets search our binary for each as they are 4-bytes long.

```
(gdb) find 0x08000000, 0xA0000000, 0xe000e100
```

Pattern not found.

Ok let's try the next one.

```
(gdb) find 0x08000000, 0xA0000000, 0xe000e104
```

1 pattern found.

Great! Let's examine what is at the address minus a few bytes.

```
(gdb) x/200i 0x08000300
```

Notice there are bytes followed by two bytes #13 and #10 this can signify a /r/n so we might have our UART message, at least one of them.

0x800044a: cmp r0, #49 ; 0x31

This is a 1 and we know from our intel this is the kill switch value.

0x80003ae: mov.w r7, #64 ; 0x40

0x80003b2: bl 0x80002d8

0x80003b6: cmp r7, #64 ; 0x40

We know from our intel that 64ms is the value we need to change to 8ms so this could be it! The first 0x40 is the actual speed and the 2nd 0x40 is the expected value.

Fire up STM32CubeProgrammer.

0x0800044a, we see D1012831 so we need to change the last byte to 01.

0x080003ae, we see the value of 0740F04F so we need to change the 40 to 08.

0x080003b6, we see the value of D1254F40 so we need to change the 40 to 08.

BLAM!